

Ad Hoc Cloud Computing From Concept to Realization

Lectures Collections

Prepared by

Prof. Fatma Omara

Ad Hoc Cloud Computing

- Current Cloud Computing is primarily based on proprietary data centers, where hundreds of thousands of dedicated servers are setup to host the cloud services.
- There are ***billions*** of ***underutilized Personal Computers*** (PCs), usually ***used*** only for a ***few hours*** per day, owned by ***individuals*** and ***organizations*** worldwide.
- The **unused compute** and **storage capacities** of the ***underutilized PCs*** owned by individuals and organizations can be **consolidated** as ***alternative*** cloud fabrics to provision broad cloud services, primarily ***infrastructure as a Service***.
- This approach referred to as "***no data center***" approach, complements the data center based cloud provision model.

Ad hoc Cloud Computing-Cont..

- The concept of *ad hoc Cloud* within an enterprise is that *harness unused resources* to *improve* overall *utilization*, *reduce energy consumption*, and *allow* organizations to take advantage of *operating their own in-house cloud*.
- An *ad hoc Cloud* harvests resources from existing **nonexclusive** available hosts used by *host users* (e.g. company employees) and exposes these resources to *cloud jobs submitted by cloud users*.
 - **Example**, aggregating unused processing and storage capacity of an University of ten thousand machines in offices and labs to represent a major untapped computing resource.

Ad Hoc Cloud Computing-Cont..

- An *ad hoc Cloud Computing* platform is developed to transform *spare resource capacity* from an infrastructure owner's locally available, but non-exclusive and unreliable infrastructure, into an *overlay Cloud platform*.
- *Cloud jobs* are submitted to *ad hoc guests*, or **VMs**, running on each of the hosts within *ad hoc Cloud*
- Applications require *extremely high levels of security*, may not be suited to the ad hoc cloud, or any Cloud implementation.
- Each host has *an ad hoc client* installed to *control* the guest, *monitor* resources and *provide state information* to the server.
- Within an ad hoc cloud, a *set of cloudlets* are created; a set of *connected ad hoc guests* that offer a particular **service** or **environment**.

Ad Hoc Computing Terminologies

Non-exclusive Infrastructure

An infrastructure whose hosts are reserved for *some primary purpose*,

- e.g., employee workstations running company applications.

Ad Hoc Host

A physical machine whose resources are donated to the *ad hoc Cloud*, but is used for some other *primary purpose*.

Host Processes

Processes which execute by Ad hoc host (i.e., *host* and *guest processes*)

Ad Hoc Computing Terminologies- Cont..

Cloudlet

- A set of connected ***ad hoc guests*** that provide a particular service or execution environment.

Note That

- Cloudlet members could be distributed across an infrastructure
- a **host** may either **dedicate** to a **single cloudlet** or **attached** to **many cloudlets** to allow applications from different domains to run concurrently upon the same host.

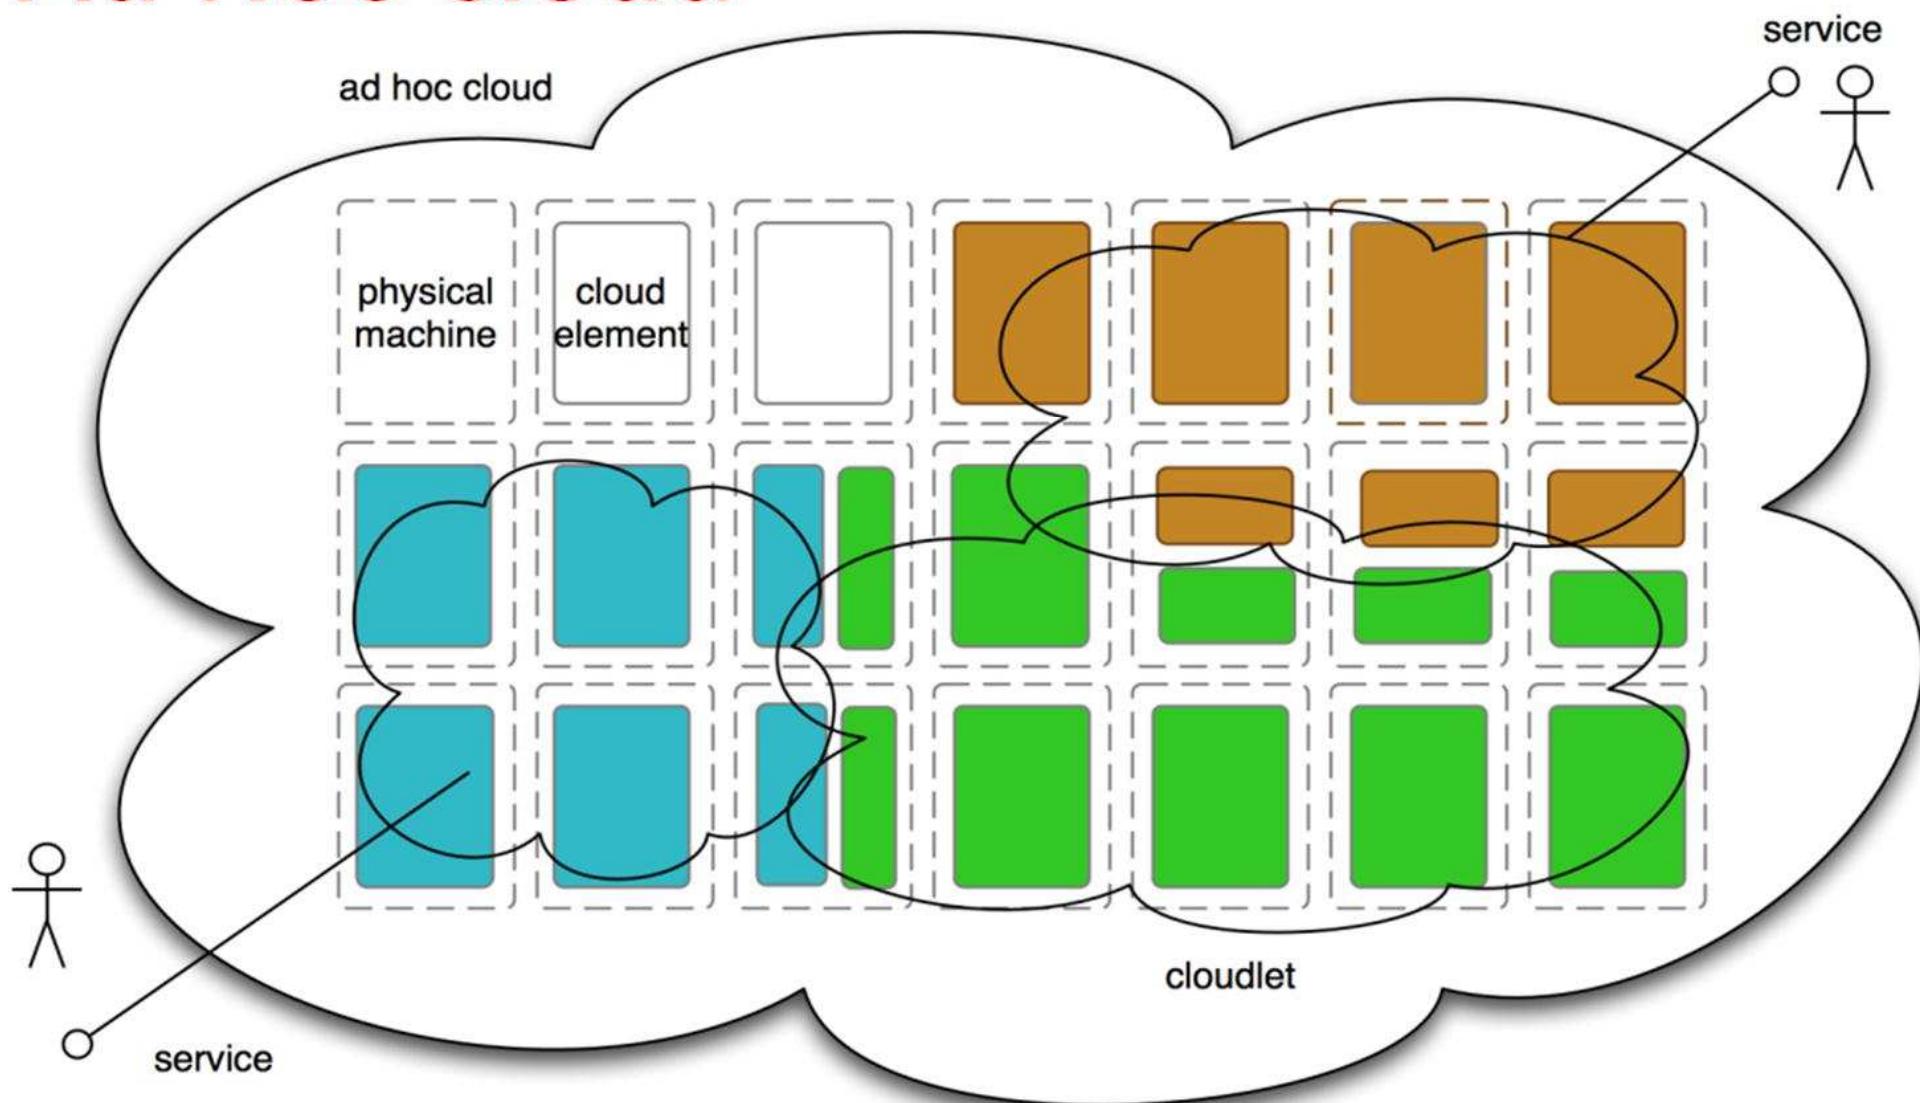
Volunteer User

A user of a volunteer computing infrastructure.

Volunteer Host

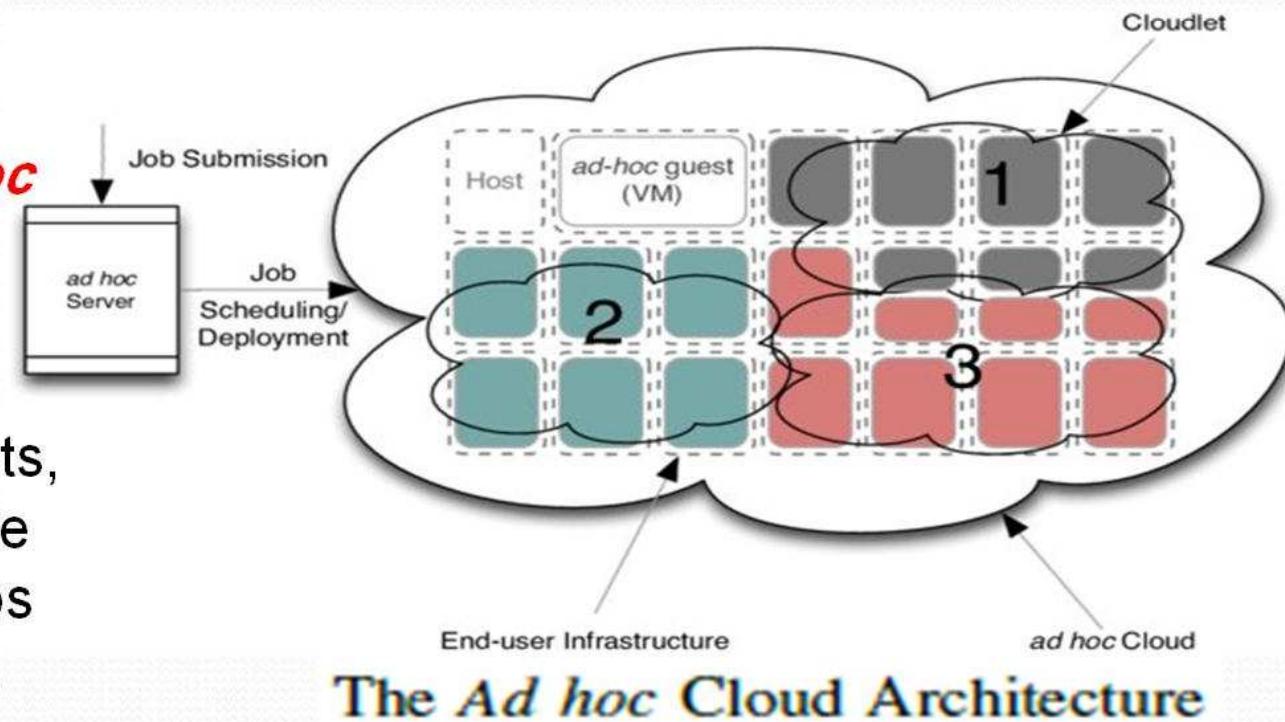
a host whose resources are donated to a volunteer computing infrastructure under the instruction of a volunteer user.

Cloudlets and Cloud Elements of an Ad Hoc Cloud



Ad hoc Cloud Computing-Cont..

- For instance, *cloudlet '1'* and *cloudlet '2'* may hold the necessary environments for *Matlab* and *BLAST* applications to execute respectively.
- *Guest* may either be dedicated to a *single cloudlet*, or part of many if *multiple jobs* are *scheduled* onto the *same guest* and *require different operating environments*.
- *Cloudlets, hosts* and *guests* are all centrally managed via the *ad hoc server*.
- The *ad hoc* server *accepts* Cloud jobs, *schedules* jobs to hosts, *maintains* system state and *ensures* cloud jobs successfully complete.



Ad hoc Cloud Processes

1. A cloud user submits a job (application and/or data) to the ad hoc cloud server via a web interface.
2. First, hosts which are instructed to *donate* their machines to ad hoc platform connects to ad hoc server. A VM is sent to hosts' machines and they await further instructions.
3. Ad hoc server schedules a job to a host for execution based on host reliability.
4. Upon a job arriving, the host starts VM and job execution begins.

Ad hoc Cloud Processes- Cont..

- 4) The installed ***ad hoc client*** software in each host dynamically ***controls VM*** resources to minimize **cloud job interference** with executing host processes. Furthermore the ***ad hoc client monitors*** the resource loads of both **host** and **guest** and ***passes*** this information to the ***ad hoc cloud server*** on a regular basis.
- 5) Periodically, ***ad hoc client*** informs ***ad hoc server*** of **host** and **VM** availability to indicate that they have not failed.
- 6) Periodically. ***ad hoc client*** captures **VM snapshots** and distributes these in a Peer-to-Peer (P2P) fashion to **other reliable and available hosts within the platform**;
 - Only the most recent snapshot is stored to minimize disk consumption.

Volunteer Computing Middleware

- Integrating *Volunteer Computing* and *virtualization* is the *backbone* of the *ad hoc cloud*.
- An open source client-server middleware platform could be used
 - Virtualized Berkeley Open Infrastructure for Network Computing (**BOINC**) as an example

References

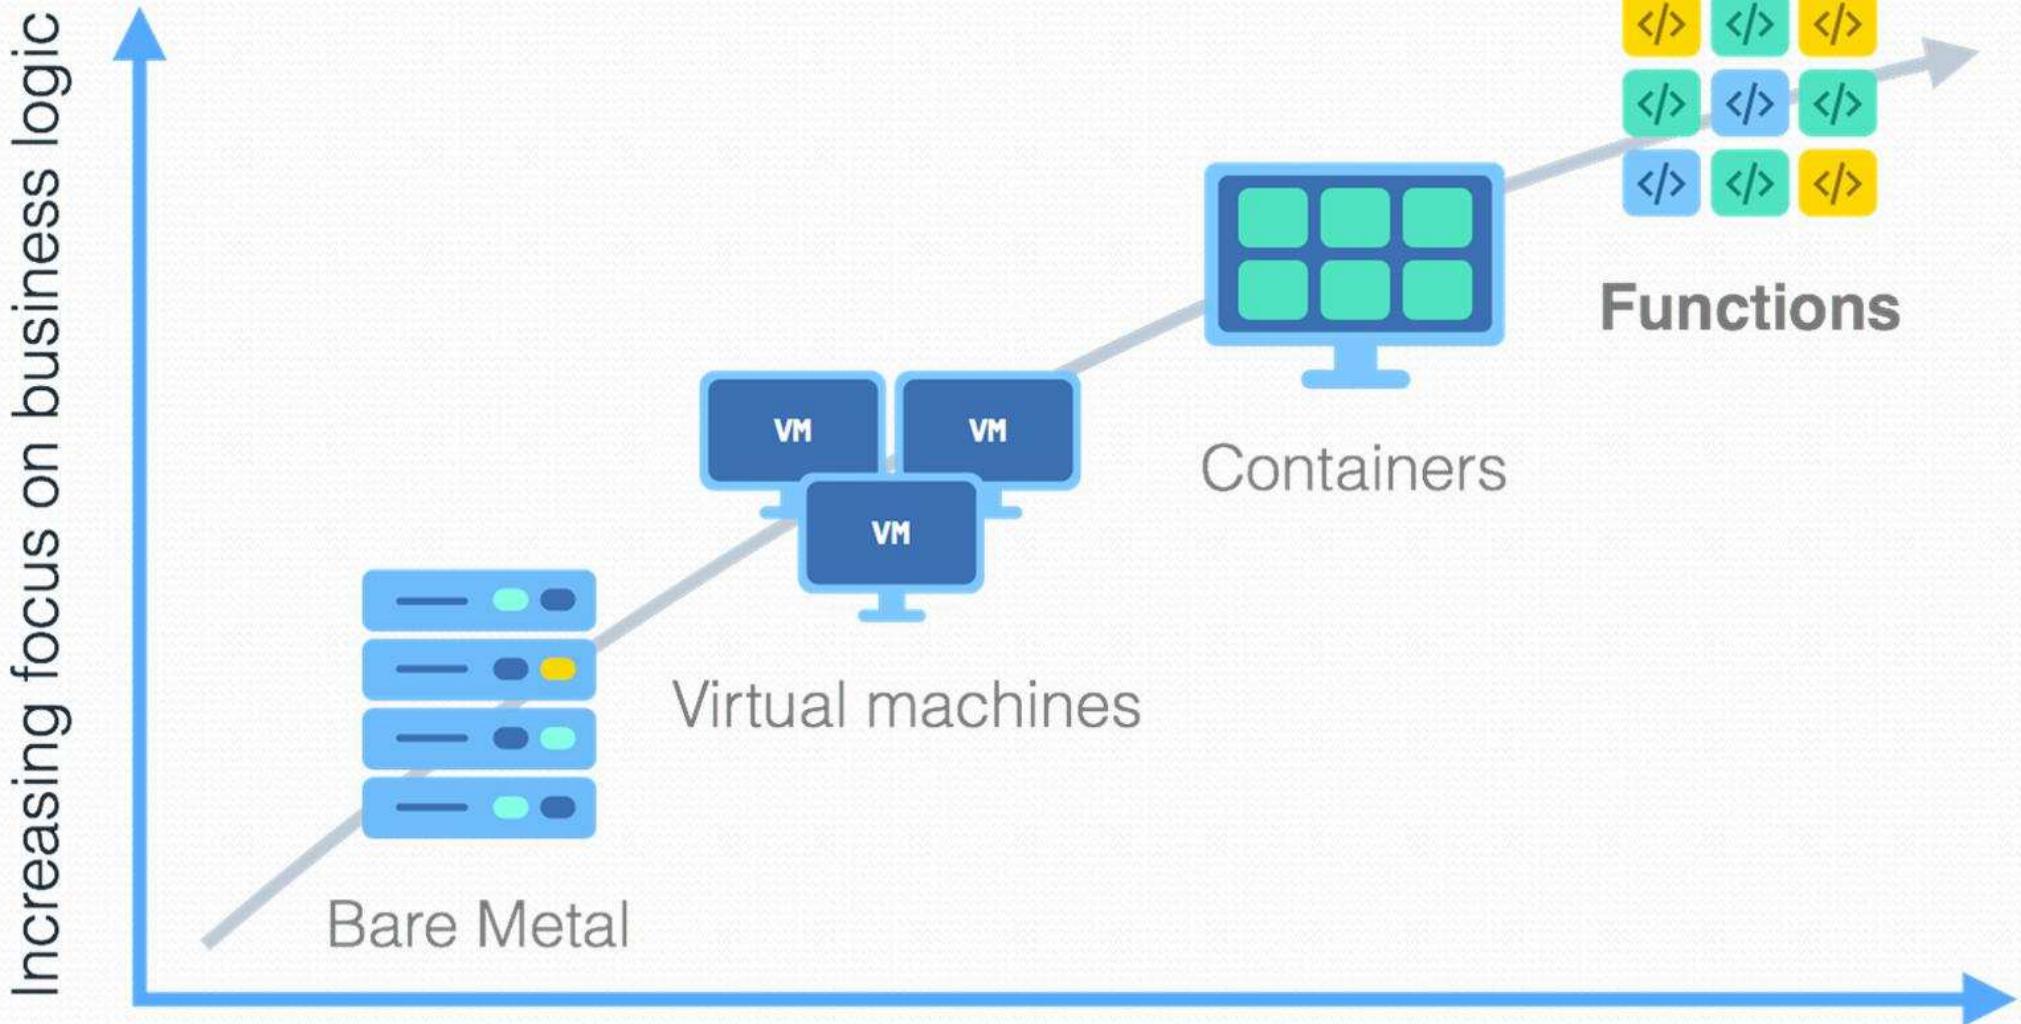
- 1) <https://arxiv.org/abs/1505.08097>
- 2) https://www.researchgate.net/publication/308868099_Ad_Hoc_Cloud_Computing
- 3) <https://ieeexplore.ieee.org/document/7214163>
- 4) http://research.nesc.ac.uk/files/Gary_McGilvary_Thesis.pdf
- 5) <https://arxiv.org/abs/1002.4738>

Serverless Computing

Prepared By: Prof. Fatma Omara

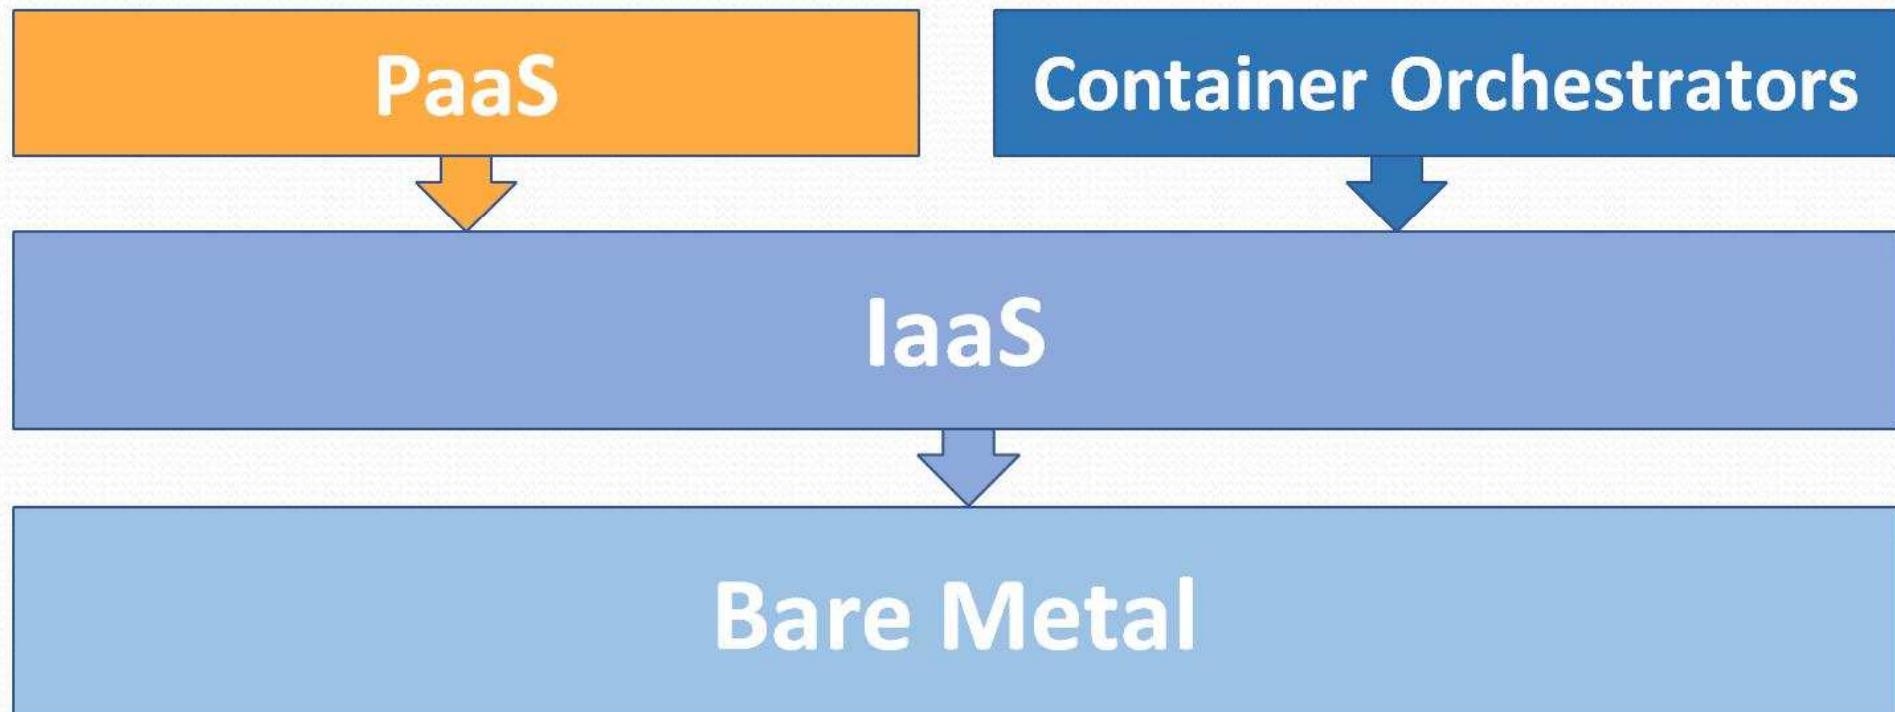
Materials from different sites and papers

Evolution of Serverless

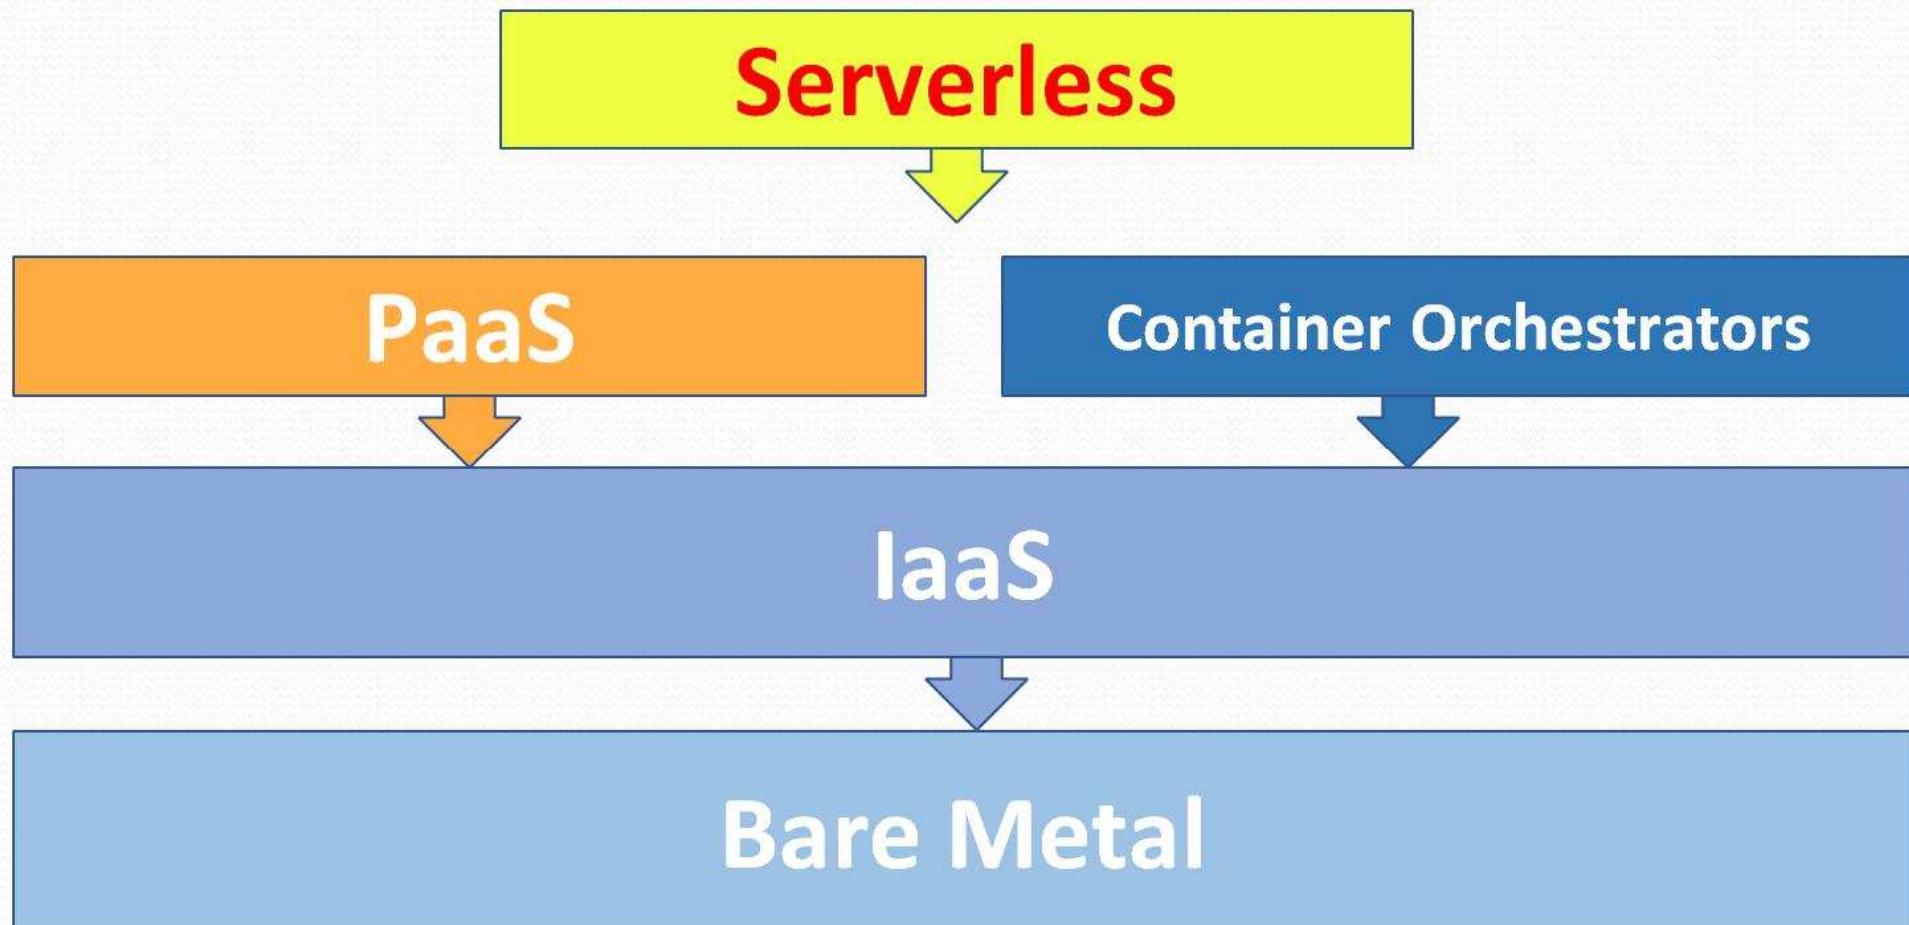


Decreasing concern (and control) over stack implementation

Evolution of Serverless



Evolution of Serverless

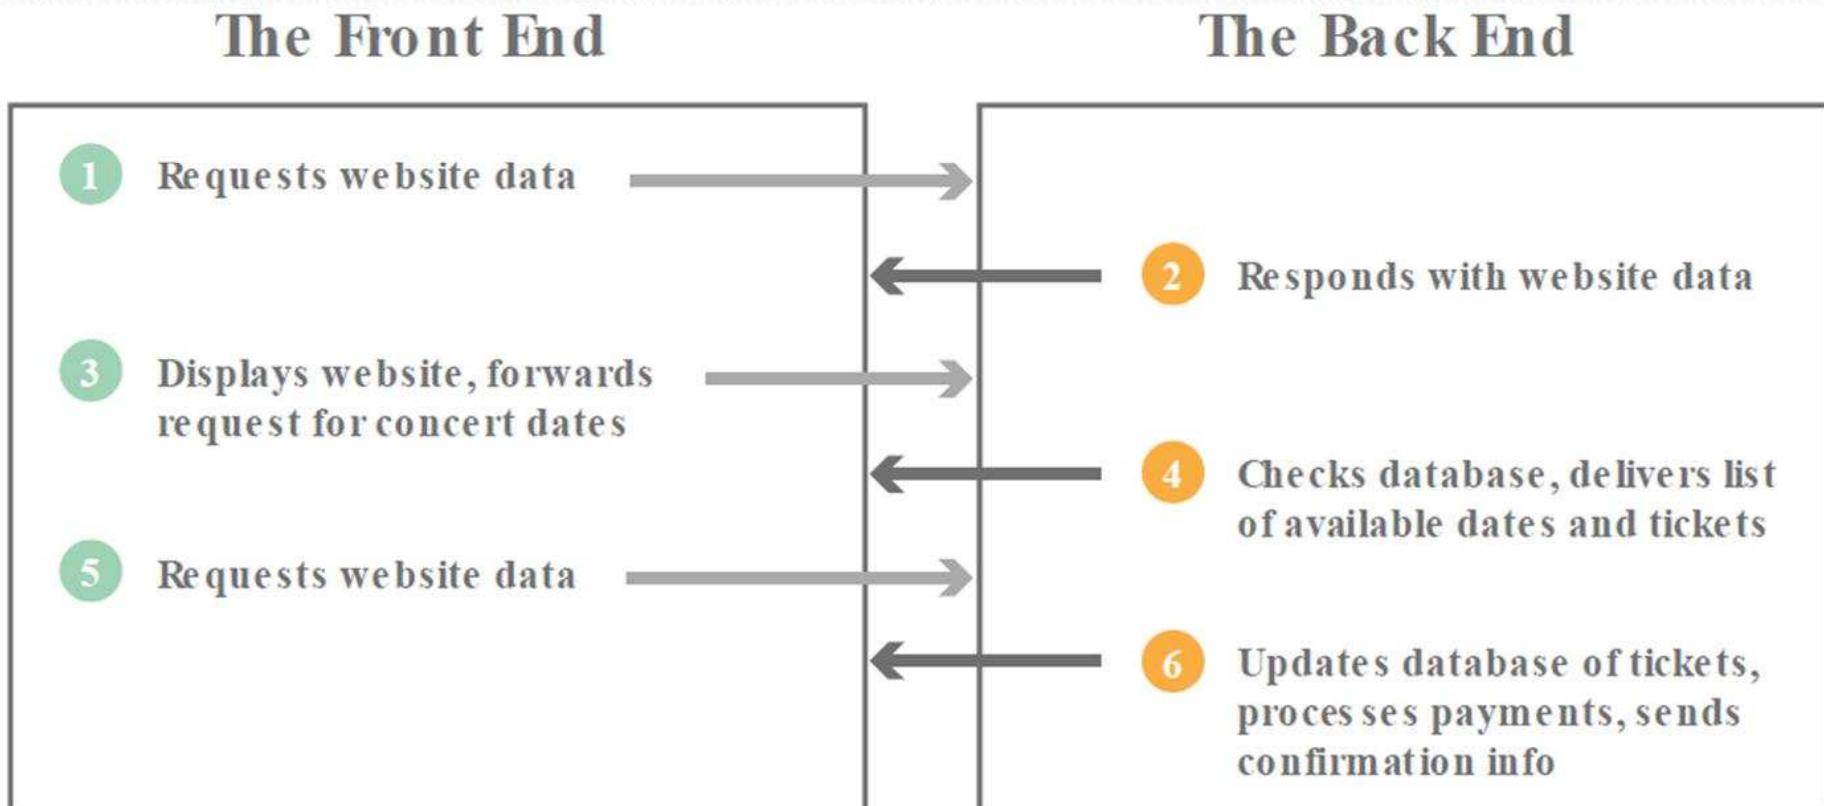


Serverless Computing

- Serverless computing is a *type* of cloud computing model, allows *users* to *write* and *deploy code*, *build web applications*, and *perform* a *range* of *other tasks*, *without* the *need* to *provision* or *manage* any *physical servers*.
- Serverless Computing **does not imply** that **servers** are not **involved**, but **servers** are utilized with administered less, i.e.,
 - **Servers operate** and **maintain** by third-party providers.
 - It **enables developers** to **deploy** and **execute** their **code** in the Cloud **without needing** to **worry** about **underlying infrastructure** they're working on, which is completely fine.
- The **service** is auto-scaling
 - No need to reserve and pay for a fixed number of servers or amount of bandwidth. Instead,
 - **Service provider** **allocates backend services** as they're required.

Serverless Computing-Cont..

- Usually, Application Development is typically split into two main sections; *frontend* and *backend*.
- **Ex:** Web Based Application



Serverless Computing-Cont..

- *Serverless computing* provides backend services using a *serverless platform*.
- *Serverless platform* is an interface through which **developers** can **build**, **deploy**, and **run applications**, **without having to worry** about the **underlying infrastructure** that's powering them.
 - *Serverless platform* offers **features** and **tools** such as Functions as a Service (FaaS), which allow for code to be triggered in response to specific, predetermined events.
- The user of a *serverless computing service* is **charged** based on **their computation only**.
- In principle, **Serverless** is similar to a cloud based phone system, where most of the associated hardware is **stored off-site** and maintained by a *third-party*.

Cloud Based Phone System

- Cloud Based Phone System is a phone service that allows you to make calls over the internet
- Cloud Phones are hosted in one or more off-site secure data centers.

Functions as a service (FaaS)

- **Function as a service (FaaS)** is a **cloud computing model** that runs code in **small modular pieces**, or **microservices**, in response to events.
- **FaaS** enables developers to **create** and **run a single function** in the **cloud** using a **serverless compute model**.
- **Cloud providers** then **manage** physical hardware, virtual machines, and web server software management.
- **FaaS** lets **developers write** and **update** a **piece** of **code** on the fly, which can be **executed** in response to an **event**, as a user **clicking** on an **element** in a web application
- Examples of FaaS are databases (SQL and NoSQL), and storage (particularly object storage)
- **Companies Offer Function-as-a-service:**
 - ❖ Alibaba Cloud Function Compute, AWS Lambda, Google Cloud Functions, IBM Cloud Functions based on Apache OpenWhisk, Microsoft Azure Functions, Oracle Cloud Functions.

How Serverless Computing Work

- In practice, Serverless computing works :
 1. Developers write code, and deploy it to their *cloud provider*.
 2. This code is then packaged by the *Cloud Provider*, and deployed to a **fleet of servers**.
 3. Upon a request being made to execute the code, the *Cloud Provider* will create a *new container* to run the code in, which is then destroyed when the execution has completed.
- Using this system, developers only need pay for the time in which their *code is executing*.

Generally,

Serverless is a cloud-native platform For:

Short-running, Stateless computation

and

event-driven applications

which

**scales up and down instantly and
automatically**

and

**charges for actual usage at a millisecond
granularity**

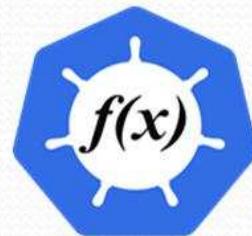
Current Platforms for Serverless



IBM Cloud Functions



Azure Functions



Red-Hat



AWS
Lambda



Google Functions



OpenLambda



Kubernetes



Advantages of Serverless Computing

Lower Costs

- Serverless Computing is charged on an event-based model.
 - Service providers charge developers for their code execution time
 - Eliminating needs of **provision, manage, and maintain physical servers**, (so, save a great deal of money).

Increase Productivity

- No need to devote time to **maintaining hardware**,
 - Allows development teams focus on actually developing.
- Server-side applications rely on functions, or series of functions, decided by provider's infrastructure, So,
 - Developers need to upload a few pieces of code and run the program.

Advantages of Serverless Computing

Greater Scalability

- Using serverless computing, developers can quickly scale their operation up or down depending on current demand.
 - If demand for functions increases, the provider's servers adjust, and provide higher capacity to run the functions.
- There's no limitation based on the storage or performance capabilities of the server.

Greener Computing

- Server less computing is considered to be a greener technology
- Resource utilization is improved because resources are only **used when they're needed to execute code**. In addition, energy saving

Serverless Computing Disadvantages

Possible Performance Issues

- Serverless computing can be *prone to performance issues* in some scenarios. If an application is reactivated after a prolonged period of non-use, it can experience a '**Cold Start**'.
 - It arises from function **starting up slowly**, and the serverless infrastructure requiring some time to process the server request. This can lead to **slower performance**.
 - To limit **cold starts occurrence**, keeping *strings of code short*
 - However, this has the adverse effect of creating a **higher number of smaller functions** to manage, which can be inconvenient.

Serverless Computing Disadvantages

Increased Complexity

- Many serverless architectures operate on a **multi-tenancy model**, so multiple different software programs for multiple different clients may be **running on servers *simultaneously***.
 - This leads to issues of ***low performance***, or ***security risks*** arising from customers being able to access one another's data.
- Serverless computing is **unsuitable** for **long-running workloads** Because
 - Serverless solutions charge based on **the amount of time that code is being run**, applications that require **lengthy processes** may be **more costly** than if they were run on dedicated servers.

Serverless Computing Disadvantages

Complicated Testing & Debugging

- The nature of serverless computing results in developers having a ***lack of backend visibility***, which can make testing and debugging quite challenging.
- Developers using serverless computing infrastructure should take **additional steps** in order to **recognize**, **predict**, and **plan** for **faults**.
 - This helps to limit the interruptions of services for users.

Serverless is *good* for

- short-running
- Stateless
- event-driven



Microservices



Mobile Backends



Bots, ML Inferencing



IoT



Modest Stream Processing



Service integration

Serverless is *not good* for

- long-running
- stateful
- number crunching



Databases



Deep Learning Training



Heavy-Duty Stream Analytics



Numerical Simulation



Video Streaming

Serverless Computing Vrs Cloud Computing

- *Platform as a service (PaaS)*, *containers* and *virtual machines (VMs)* all play a **critical role** in the cloud application development and compute ecosystem
- **Serverless** compares to **Cloud** across some **key attributes**:
 - **Provisioning Time:** Measured in *milliseconds* for serverless, vs. *minutes* to *hours* for the other models.
 - **Administrative burden:** *None* for serverless, compared to a continuum from light to medium to heavy for PaaS, containers and VMs respectively.
 - **Maintenance:** Serverless architectures are *managed 100%* by the *provider*. The *same* is *true* for **PaaS**, but *containers* and **VMs** require significant maintenance including updating and managing operating systems, container images, connections, and so on.

Serverless and Microservices

- Given its unique *combination* of attributes and benefits, **serverless** architecture *works best* on *use cases* around **microservices**, **mobile backends** and **data** and **event stream processing**.
- The **most common use case** of **serverless** today is **supporting microservices architectures**.
- Serverless gained significant momentum given its attributes around **small bits of code**, **inherent automatic scaling**, **rapid provisioning** and a **pricing model** that **never charges** for **idle capacity**.

References

1. <https://www.codemotion.com/magazine/devops/cloud/serverless-computing-the-advantages-and-disadvantages/>
2. <https://www.ibm.com/topics/serverless>
3. <https://www.cloudflare.com/learning/serverless/what-is-serverless/>
4. [https://www.google.com/search?q=functions+as+a+services+as+a+service+\(FaaS\)&rlz=1C1GGRV_enEG751EG755&oq=functions+as+a+functions+as+a+service+\(FaaS\)+&aqs=chrome..69i57j33i1o16ol5.13888joj7&sourceid=chrome&ie=UTF-8](https://www.google.com/search?q=functions+as+a+services+as+a+service+(FaaS)&rlz=1C1GGRV_enEG751EG755&oq=functions+as+a+functions+as+a+service+(FaaS)+&aqs=chrome..69i57j33i1o16ol5.13888joj7&sourceid=chrome&ie=UTF-8)