

NLP Lab -6-

Using Word Embedding and CNN to Build a Classifier

*Problem: classify news articles into pre-defined 20 categories

- Given a dataset of news article, the objective is to build a CNN classifier to predict the class of a given news article, Glove word embedding model will be used to numerically represent the articles' text
- This script loads pre-trained word embeddings (GloVe embeddings) into a frozen Keras Embedding layer, and uses it to train a text classification model on the 20 Newsgroup dataset
- GloVe embedding data can be found at: <http://nlp.stanford.edu/data/glove.6B.zip> (source page: <http://nlp.stanford.edu/projects/glove/>)
- 20 Newsgroup data can be found at: <http://www.cs.cmu.edu/afs/cs.cmu.edu/project/theo-20/www/data/news20.html> (<http://www.cs.cmu.edu/afs/cs.cmu.edu/project/theo-20/www/data/news20.html>)
- reference: <https://blog.keras.io/using-pre-trained-word-embeddings-in-a-keras-model.html>

(1) Imports

In [2]:

```
import os
import sys
import numpy as np
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
from keras.utils import to_categorical
from keras.layers import Embedding
from keras.layers import Dense, Input, GlobalMaxPooling1D
from keras.layers import Conv1D, MaxPooling1D, Embedding, Flatten
from keras.models import Model
```

```
c:\users\sarahhassan\appdata\local\programs\python\python35\lib\site-packages\h5py\__init__.py:36: FutureWarning: Conversion of the second argument of issubdtype from `float` to `np.floating` is deprecated. In future, it will be treated as `np.float64 == np.dtype(float).type`.
```

```
from ._conv import register_converters as _register_converters
Using TensorFlow backend.
```

```
c:\users\sarahhassan\appdata\local\programs\python\python35\lib\site-packages\requests\__init__.py:80: RequestsDependencyWarning: urllib3 (1.22) or chardet (2.3.0) doesn't match a supported version!
```

```
RequestsDependencyWarning)
```

(2) Load dataset

prepare text samples and their labels.

- Input: path to a folder containing news files. the dataset folder contain 20 folder each representing one class. each folder contains samples of news articles from each class. each article is written in a separate text file.
- output:

1) list of all text samples

2) list of the label ID for each text sample

3) dictionary mapping label ID to label name

In [3]:

```
print('Processing text dataset')

TEXT_DATA_DIR = '20_newsgroup'
texts = [] # list of text samples
labels_index = {} # dictionary mapping label name to numeric id
labels = [] # list of label ids
for name in sorted(os.listdir(TEXT_DATA_DIR)):
    path = os.path.join(TEXT_DATA_DIR, name)
    if os.path.isdir(path):
        label_id = len(labels_index)
        labels_index[name] = label_id
        for fname in sorted(os.listdir(path)):
            if fname.isdigit():
                fpath = os.path.join(path, fname)
                args = {} if sys.version_info < (3,) else {'encoding': 'latin-1'}
                with open(fpath, **args) as f:
                    t = f.read()
                    i = t.find('\n\n') # skip header
                    if 0 < i:
                        t = t[i:]
                    texts.append(t)
                    labels.append(label_id)
```

Processing text dataset

In [4]:

```
print('Found %s texts.' % len(texts))
print('nLabels = ', len(labels))
print('Classes are:\n ')
for key in labels_index:
    print (key)
```

Found 19997 texts.

nLabels = 19997

Classes are:

```
alt.atheism
soc.religion.christian
comp.windows.x
sci.crypt
comp.sys.ibm.pc.hardware
sci.med
comp.os.ms-windows.misc
rec.motorcycles
comp.sys.mac.hardware
rec.sport.hockey
talk.religion.misc
comp.graphics
misc.forsale
sci.electronics
sci.space
talk.politics.mideast
rec.sport.baseball
rec.autos
talk.politics.guns
talk.politics.misc
```

(3) Process Text samples

The steps to process text samples include:

- tokenize samples to words and keep top 20,000 most commonly occurring words in the dataset and neglect the others
- form a dictionary of all words in the dataset.. the dictionary maps a word to a word ID
- replace each word by its ID
- truncate/pad the sequences to a maximum length of 1000 words

** Keras tokenizer object lowers all characters and removes special characters: '!"#\$%&()+,-./:;<=>?*

@[]^_`{ }~\t\n'

=> Each news article will be a vector of words IDS

In [5]:

```
MAX_SEQUENCE_LENGTH = 1000
MAX_NUM_WORDS = 20000
```

In [7]:

```
tokenizer = Tokenizer(num_words=MAX_NUM_WORDS)
tokenizer.fit_on_texts(texts)
sequences = tokenizer.texts_to_sequences(texts)

word_index = tokenizer.word_index # the dictionary
print('Found %s unique tokens.' % len(word_index)) #only top MAX_NUM_WORDS will be used
to generate the sequences
data = pad_sequences(sequences, maxlen=MAX_SEQUENCE_LENGTH)
print('Shape of samples:', data.shape)
print('Sampele:(the zeros at the begining are for padding text to max length)')
print(data[2])
```

Found 174074 unique tokens.

Shape of samples: (19997, 1000)

Sampele:(the zeros at the begining are for padding text to max length)

```
[ 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 6 54
10754 1731 26 8090 201 1731 26 9652 8455 43 107 316
 52 4 346 251 17 1341 68 4299 39 8 58 362
 9 51 200 18 575 3020 5 9 1 2694 3 89
 25 575 2 166 24 17 6 4 648 14 8 4
108 610 2 57 105 13 6174 4 246 1 610 654
 21 1878 89 2518 1304 6 316 24 17 6 1 521
120 27 70 6 1456 22 3020 240 27 19 817 2
 16 597 262 2 180 2 883 3 89 21 135 5
 17 32 1629 251 432 1507 9 51 1183 3020 38 1032
265 166 154 3020 5 38 252 362 12 1 2201 2400
 7 65 84 14 1878 23 13 179 1 134 3231 91
3020 182 18 575 316 24 316 8 32 1056 6 282
238 713 599 3 93 146 61 19 107 567 1 1259
 3 18377 46 19 599 799 3 4089 146 17 1 1190
 3 1 3187 1 2106 2288 12 1 93 329 3 429
1456 5 3020 88 17 559 71 9483 5626 11 289 1456
 52 205 1 666 3 429 5 3020 9 3 1456 5
255 9 3 429 21 11 8 2860 9 316 1183 1
2694 3 15003 666 1 362 12 9 8 17 6175 7221
 28 1351 11222 3 316 18 83 2270 692 142 5 61
 80 19 27 8593 80 309 19 27 1140 7 690 216
 1 414 31 501 11 25 7 84 6 1 1185 15004
 22 729 5 11 25 1370 21 309 21 9190 1015 67
 27 19 29 44 1015 121 42 27 7403 834 15 1
329 15 1677 316 154 1 687 3 1 342 7 65
116 53 180 11 8 615 35 802 145 27 81 50
126 501 883 3 1 1143 5 11 1264 9 31 258
123 3 77 151 1950 4089 21 1 1019 3 71 1295
 5 114 1 2366 3 3020 2 543 22 429 2 870
 19 17 1202 1687 107 4 6764 1231 3 870 38 16
849 262 513 11 203 32 485 861 20 1495 344 3
```

```

5562      5 4194      38      16      637      48      827      9 1025      2      130
112       7 1859      159      5       7      65      84      4      40 1491 2604
432       8      9      349      23 3020      5      429      5 1456 2422      71
2309 1043      29      91      7      157      419      163      6      1      576      3
397 10568      29      320      2 3189      9      58      66 15005      6      256
1168      40 1202      773      50      184      185      82      14      5      23      316
8         4      91      452      290      68      2      16      145      9      316      25
4 13456      8      17 1202 1687      1 1922      5 1031      861      20
1 4299      19      661      151      21      1      610      9      316      25      4
13456 12815      15      1      476      6      1 3189      3      316      667      8
39       55      75      362      12      11      40 1491      5      40 2031 2604
8        349      114      6      138      387 2051      9      2 1168      3      597
1143      6      138 1719      6      521 13457 6762      17      114      2      840
174       9 4570      19      17 1341      1      259      835      11      8      86
1264      9      429      52      83 5766      80      309      19      1 10755 11222
2         33      883      61      34      16      1 10755      8      346      154      55
3        100 6680      5      11      8      17      114      945      1      76      580
5715      8      168      3      1 6176      5      96 1350      4 19951 4472
1 10755 1680      252      4      858      579      3      2      1      180      142
1 17068 1680      391      2 1111      174      6      1      134 1264 1255
8         9      31      58      180      230      3      1      666      25      750      1
1307      29 3083      128      15      8      457      917      538      5      253      2
3105      4      3      1      305      3 3020      6      75      560      40      88
17       74      35      1      485      3      429      139      221      64      5 1403
323      15      429      19      457 2298      24      80      8      9 2069      2
4         3      316 5816]

```

(4) format output of the CNN (the shape of labels)

- Each label will be a binary vector of size 20

In [8]:

```

labels_matrix = to_categorical(np.asarray(labels))
print('Shape of data tensor:', data.shape)
print('Shape of label tensor:', labels_matrix.shape)
print('Sample label:\n', labels_matrix[1590])

```

Shape of data tensor: (19997, 1000)

Shape of label tensor: (19997, 20)

Sample label:

```
[0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
```

(5) Split samples and labels to training and testing sets

In [9]:

```
VALIDATION_SPLIT = 0.2

indices = np.arange(data.shape[0])
np.random.shuffle(indices)
data_shuffled = data[indices]
labels_shuffled = labels_matrix[indices]
nb_validation_samples = int(VALIDATION_SPLIT * data_shuffled.shape[0])

x_train = data_shuffled[:-nb_validation_samples]
y_train = labels_shuffled[:-nb_validation_samples]
x_val = data_shuffled[-nb_validation_samples:]
y_val = labels_shuffled[-nb_validation_samples:]
print('Shape of training data: ',x_train.shape)
print('Shape of testing data: ',x_val.shape)
```

Shape of training data: (15998, 1000)

Shape of testing data: (3999, 1000)

(6) Read Glove Word Embeddings

Build a dictionary mapping words in the embeddings set to their embedding vector

In [10]:

```
EMBEDDING_DIM = 100

print('Indexing word vectors.')

embeddings_index = {}

with open('glove.6B.100d.txt') as f:
    for line in f:
        values = line.split(sep=' ')
        word = values[0]
        coefs = np.asarray(values[1:], dtype='float32')
        embeddings_index[word] = coefs

print('Found %s word vectors.' % len(embeddings_index))
```

Indexing word vectors.

Found 400000 word vectors.

(7) Map the dataset dictionary of (words,IDs) to a matrix of the embeddings of each word in the dictionary

the words in the dataset that don't exist in Glove's dictionary will get a zeroes vector these words already have id 0 so there zeros embedding vector will be at index zero in the new matrix

In [11]:

```
embedding_matrix = np.zeros((len(word_index) + 1, EMBEDDING_DIM))#+1 to include the zeros vector for non-existing words
for word, i in word_index.items():
    embedding_vector = embeddings_index.get(word)
    if embedding_vector is not None:
        # words not found in embedding index will be all-zeros.
        embedding_matrix[i] = embedding_vector
print ('Shape of Embedding Matrix: ',embedding_matrix.shape)
```

Shape of Embedding Matrix: (174075, 100)

(8) Build the Deep NN:

(8.1) Embedding Layer

responsible of converting a padded sequence of words IDs to a sequence of words embeddings

In [12]:

```
embedding_layer = Embedding(len(word_index) + 1, #vocab size
                             EMBEDDING_DIM, #embedding vector size
                             weights=[embedding_matrix], #weights matrix
                             input_length=MAX_SEQUENCE_LENGTH, #padded sequence length
                             trainable=False)
```

(8.2) Build 1D CNN Layers

In [13]:

```
sequence_input = Input(shape=(MAX_SEQUENCE_LENGTH,), dtype='int32')
embedded_sequences = embedding_layer(sequence_input)
x = Conv1D(128, 5, activation='relu')(embedded_sequences)
x = MaxPooling1D(5)(x)
x = Conv1D(128, 5, activation='relu')(x)
x = MaxPooling1D(5)(x)
x = Conv1D(128, 5, activation='relu')(x)
x = MaxPooling1D(35)(x) # global max pooling
x = Flatten()(x)
x = Dense(128, activation='relu')(x)
preds = Dense(len(labels_index), activation='softmax')(x)
```

(8.3) Build, Compile, and Run the model

In [14]:

```
model = Model(sequence_input, preds)
model.compile(loss='categorical_crossentropy',
              optimizer='rmsprop',
              metrics=['acc'])

# happy learning!
model.fit(x_train, y_train, validation_data=(x_val, y_val),
          epochs=5, batch_size=128)
```

Train on 15998 samples, validate on 3999 samples

Epoch 1/5

15998/15998 [=====] - 189s 12ms/step - loss: 2.43
14 - acc: 0.2108 - val_loss: 1.8208 - val_acc: 0.3576

Epoch 2/5

15998/15998 [=====] - 183s 11ms/step - loss: 1.54
58 - acc: 0.4616 - val_loss: 1.3524 - val_acc: 0.5279

Epoch 3/5

15998/15998 [=====] - 183s 11ms/step - loss: 1.19
46 - acc: 0.5871 - val_loss: 1.1309 - val_acc: 0.6107

Epoch 4/5

15998/15998 [=====] - 182s 11ms/step - loss: 0.97
26 - acc: 0.6671 - val_loss: 1.0934 - val_acc: 0.6297

Epoch 5/5

15998/15998 [=====] - 182s 11ms/step - loss: 0.83
66 - acc: 0.7165 - val_loss: 0.9568 - val_acc: 0.6734

Out[14]:

<keras.callbacks.History at 0x375c8a20>

(8.4) Evaluate the model

In [15]:

```
print('Acuracy on testing set:')
model.evaluate(x_val,y_val)
```

Acuracy on testing set:

3999/3999 [=====] - 18s 5ms/step

Out[15]:

[0.9568204918066303, 0.6734183546631716]

(9) Use the model for prediction

In [16]:

```
model.predict(x_val)
```

Out[16]:

```
array([[1.0327732e-02, 2.9313486e-04, 2.4011679e-06, ..., 1.7724369e-05,
        1.8570792e-03, 7.6327776e-04],
       [1.2563624e-02, 1.9062838e-02, 8.7171635e-03, ..., 8.0603240e-03,
        2.9240904e-02, 5.6615467e-03],
       [5.5559494e-06, 4.2651324e-03, 4.3965750e-03, ..., 2.5074318e-04,
        1.2648354e-03, 9.5650757e-06],
       ...,
       [1.1672964e-03, 2.3448658e-05, 2.9504277e-05, ..., 1.2243649e-03,
        3.6598616e-03, 3.0622948e-03],
       [4.8667334e-06, 7.7692151e-02, 5.7424837e-01, ..., 6.6595958e-06,
        6.7728670e-06, 1.7917351e-05],
       [2.5362926e-06, 2.0805377e-01, 3.7774213e-02, ..., 6.3202569e-06,
        1.6121081e-05, 2.8711088e-06]], dtype=float32)
```

In [18]:

```
sample = 1
label_vec = model.predict(data[sample].reshape(1,-1))
label_id = np.argmax(label_vec)
label_name = ''
for name, ID in labels_index.items():    # for name, age in dictionary.iteritems(): (f
or Python 2.x)
    if label_id == ID:
        label_name = name
        break
print ('The category of article no %s is %s' %(sample ,label_name))
```

The category of article no 1 is soc.religion.christian