

# Information Retrieval

Instructor : Dr. Hanaa Bayomi Ali

Mail : h.mobarz @ fci-cu.edu.eg

# Information Retrieval

Information Retrieval (IR) is finding material (usually documents) of an unstructured nature (usually text) that satisfies an information need from within large collections(usually stored on computers).

# Information Retrieval

- What do we understand by **documents**? How do we decide what is a document and whatnot?
- What is an **information need**? What types of information needs can we satisfy automatically?
- What is a **large collection**? Which environments are suitable for IR

# Basic assumptions of Information Retrieval

- **Collection**: A set of documents
  - Assume it is a static collection
- **Goal**: Retrieve documents with information that is **relevant** to the user's **information need** and helps the user complete a **task**

# Why IR is hard?

IR is mostly about *relevance*

- Relevance is the core concept in IR, but nobody has a good definition
- Relevance = useful
- Relevance = topically related
- Relevance = new
- Relevance = interesting
- Relevance = ???
- However we still want *relevant information*

# Why IR is hard?

- Information needs must be expressed as a query
  - But users don't often know what they want
- Problems
  - Verbalizing information needs
  - Understanding query syntax
  - Understanding search engines

# Why this is hard?

- Documents/images/ video/speech/etc are complex. We need some representation
- Semantics
  - What do words mean?
- Natural language
  - How do we say things?
- ☹ Computers cannot deal with these easily

# Semantics

Bank Note



River Bank



Bank



Blood bank





# Information Retrieval Techniques

- Index Terms (Attribute) Selection:
  - Stop list
  - Word stem
  - Index terms weighting methods
- Terms **X** Documents Frequency Matrices
- Information Retrieval Models:
  - Boolean Model (incident matrix)
  - Vector Model
  - Probabilistic Model

# Boolean Model

- Consider that index terms are either present or absent in a document
- As a result, the index term weights are assumed to be all binaries
- A query is composed of index terms linked by three connectives: **not**, **and**, and **or**
  - e.g.: car **and** repair, plane **or** airplane
- The Boolean model predicts that each document is either relevant or non-relevant based on the match of a document to the query

# Example

*Brutus AND Caesar BUT NOT Calpurnia*

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0



1 if **play** contains  
**word**, 0 otherwise

# Example

So we have a 0/1 vector for each term.

To answer query: take the vectors for *Brutus*, *Caesar* and *Calpurnia* (complemented) → bitwise *AND*.

- 110100 *AND*
- 110111 *AND*
- 101111 =
- **100100**

	 Antony and Cleopatra	Julius Caesar	The Tempest	 Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0

# Boolean Models – Problems

- Very rigid: AND means all; OR means any.
- Difficult to express complex user requests.
- Difficult to control the number of documents retrieved.
  - *All* matched documents will be returned.
- Difficult to rank output.
  - *All* matched documents logically satisfy the query.
- Difficult to perform relevance feedback.
  - If a document is identified by the user as relevant or irrelevant, how should the query be modified?

# Indexing Techniques


- Inverted index
  - Maintains two hash- or B+-tree indexed tables:
    - **document\_table**: a set of document records <doc\_id, postings\_list>
    - **term\_table**: a set of term records, <term, postings\_list>
  - Answer query: Find all docs associated with one or a set of terms
  - + easy to implement
  - – do not handle well synonymy and polysemy, and posting lists could be too long (storage could be very large)

# Example

For each term  $t$ , we must store a list of all documents that contain  $t$ .

- Identify each doc by a **docID**, a document serial number

Can we use fixed-size arrays for this?

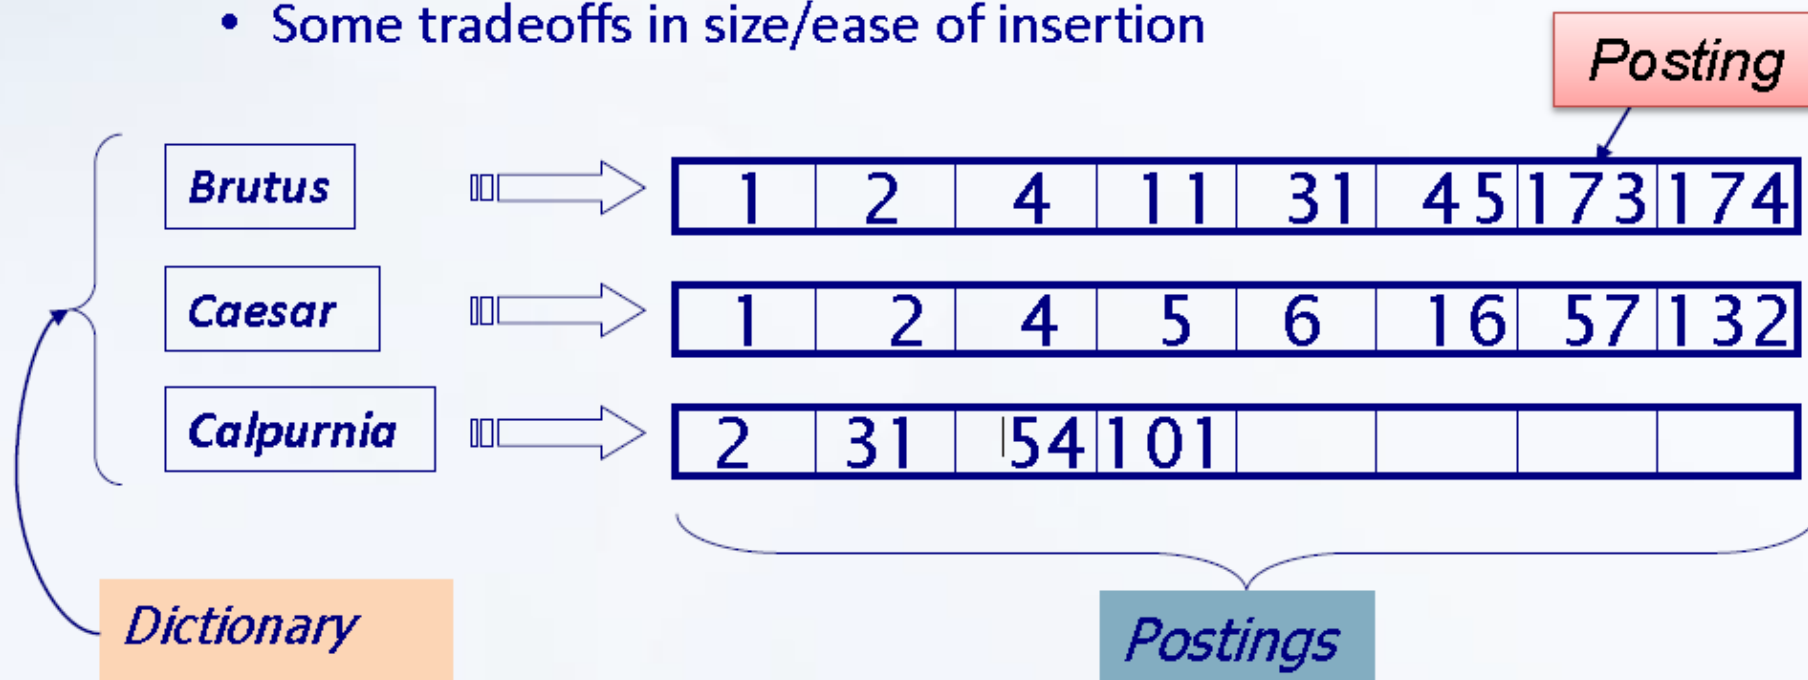


What happens if the word **Caesar** is added to document 14?

# Example

We need variable-size **postings** lists

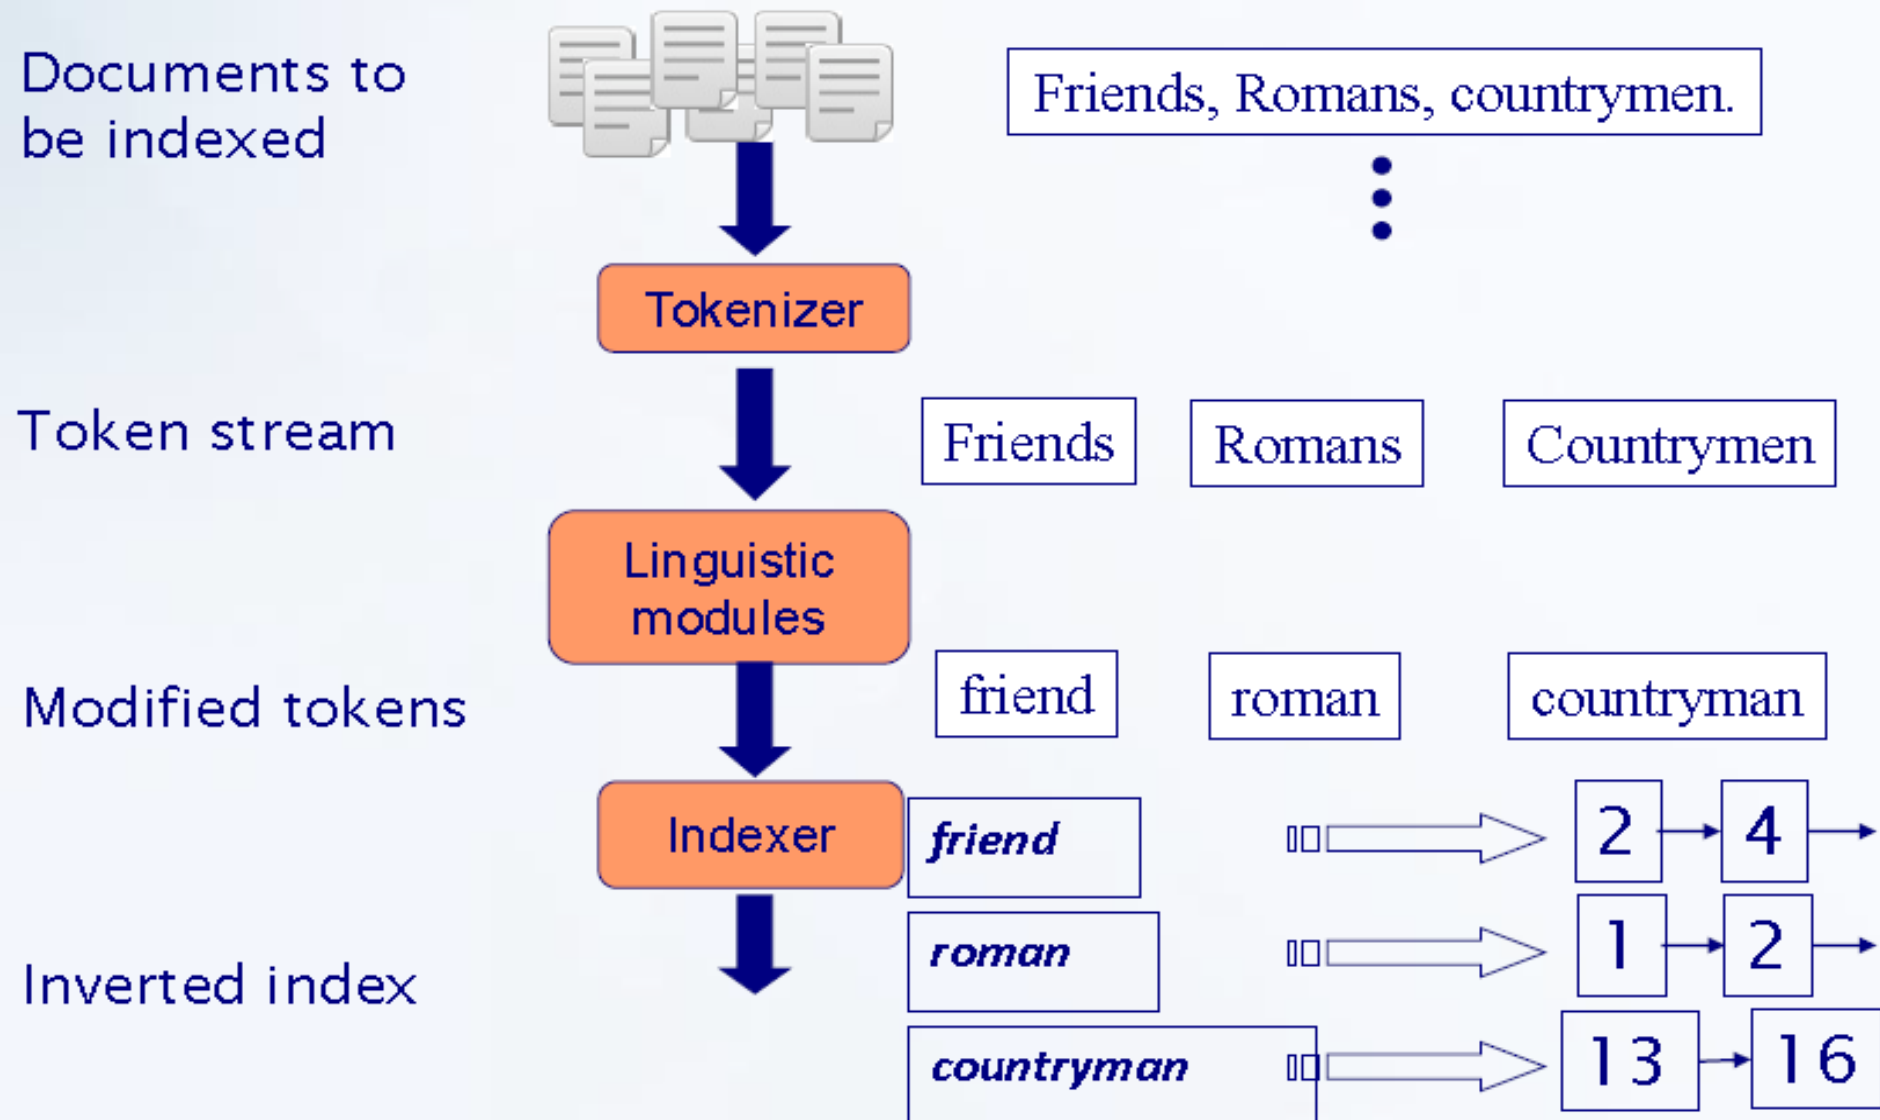
- On disk, a continuous run of postings is normal and best
- In memory, can use linked lists or variable length arrays
  - Some tradeoffs in size/ease of insertion



Sorted by docID (more later on why).



# Inverted index construction



# Initial Stages of Text Processing

## Tokenization

- Cut character sequence into word tokens (=what is a word!)
  - Deal with *“John’s”, a state-of-the-art solution*

## Normalization

- Map text and query term to same form
  - You want *U.S.A.* and *USA* to match

## Stemming

- We may wish have different forms of a root to match
  - *authorize, authorization*

## Stop words

- We may omit very common words (or not)
  - *the, a, to, of, over, between, his, him*

# Indexer Steps: Token Sequence

Sequence of (Modified token, Document ID) pairs.

Doc 1

I did enact Julius  
Caesar I was killed  
i' the Capitol;  
Brutus killed me.

Doc 2

So let it be with  
Caesar. The noble  
Brutus hath told you  
Caesar was ambitious



Term	docID
I	1
did	1
enact	1
julius	1
caesar	1
I	1
was	1
killed	1
i'	1
the	1
capitol	1
brutus	1
killed	1
me	1
so	2
let	2
it	2
be	2
with	2
caesar	2
the	2
noble	2
brutus	2
hath	2
told	2
you	2
caesar	2
was	2
ambitious	2

# Indexer Steps: Sort

Sort by terms

– And then docID

Core indexing step

Term	docID
I	1
did	1
enact	1
julius	1
caesar	1
I	1
was	1
killed	1
i'	1
the	1
capitol	1
brutus	1
killed	1
me	1
so	2
let	2
it	2
be	2
with	2
caesar	2
the	2
noble	2
brutus	2
hath	2
told	2
you	2
caesar	2
was	2
ambitious	2



Term	docID
ambitious	2
be	2
brutus	1
brutus	2
capitol	1
caesar	1
caesar	2
caesar	2
did	1
enact	1
hath	1
I	1
I	1
i'	1
it	2
julius	1
killed	1
killed	1
let	2
me	1
noble	2
so	2
the	1
the	2
told	2
you	2
was	1
was	2
with	2

# Indexer Steps: Dictionary & Postings

Multiple term entries in a single document are merged.

Split into Dictionary and Postings

Doc. frequency information is added.

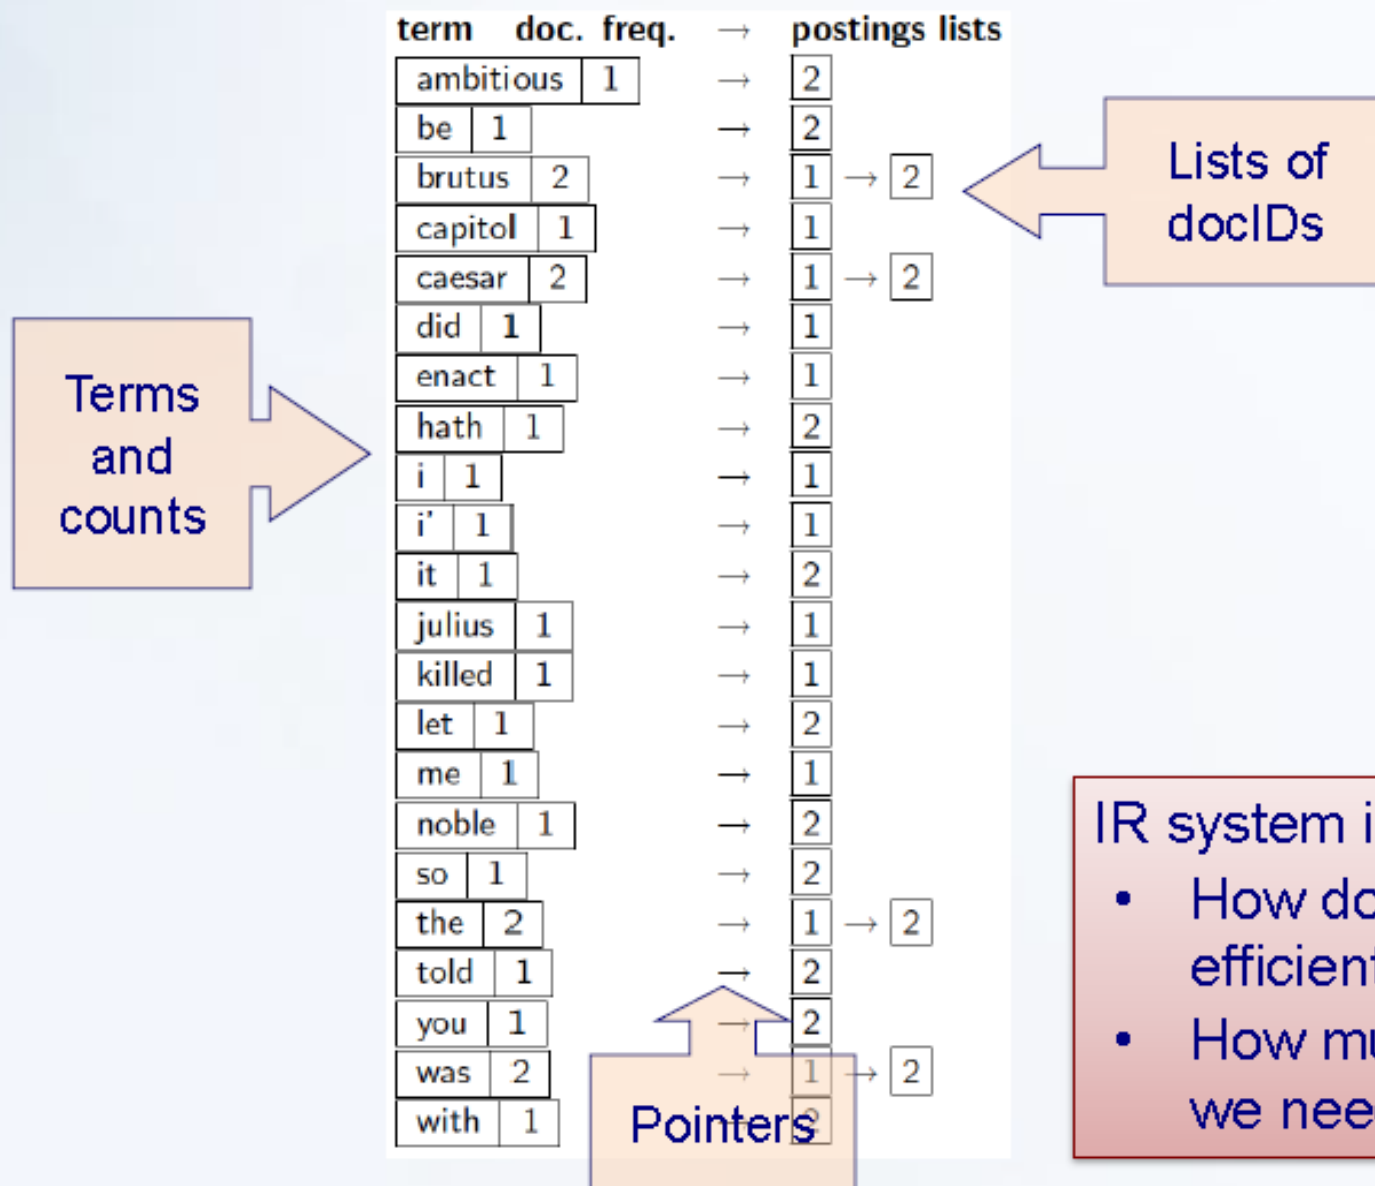
Why frequency?

Term	docID
ambitious	2
be	2
brutus	1
brutus	2
capitol	1
caesar	1
caesar	2
caesar	2
did	1
enact	1
hath	1
I	1
I	1
i'	1
it	2
julius	1
killed	1
killed	1
let	2
me	1
noble	2
so	2
the	1
the	2
told	2
you	2
was	1
was	2
with	2



term	doc. freq.	→	postings lists
ambitious	1	→	2
be	1	→	2
brutus	2	→	1 → 2
capitol	1	→	1
caesar	2	→	1 → 2
did	1	→	1
enact	1	→	1
hath	1	→	2
i	1	→	1
i'	1	→	1
it	1	→	2
julius	1	→	1
killed	1	→	1
let	1	→	2
me	1	→	1
noble	1	→	2
so	1	→	2
the	2	→	1 → 2
told	1	→	2
you	1	→	2
was	2	→	1 → 2
with	1	→	2

# Where do we pay in storage?



## IR system implementation


- How do we index efficiently?
- How much storage do we need?

## Query processing: AND

Consider processing the query:


***Brutus AND Caesar***

- Locate ***Brutus*** in the Dictionary;
  - Retrieve its postings.
- Locate ***Caesar*** in the Dictionary;
  - Retrieve its postings.
- “Merge” the two postings (intersect the document sets):



## The Merge

Walk through the two postings simultaneously, in time linear in the total number of postings entries






# Vector Space Model

- Represent a doc by a term vector
  - Term: basic concept, e.g., word or phrase
  - Each term defines one dimension
  - N terms define a N-dimensional space
  - Element of vector corresponds to term weight
  - E.g.,  $d = (x_1, \dots, x_N)$ ,  $x_i$  is “importance” of term  $i$
- New document is assigned to the most likely category based on vector similarity.

# VS Model: Illustration



# What VS Model Does Not Specify

- **How to select terms to capture “basic concepts”**
  - Word stopping
    - e.g. “a”, “the”, “always”, “along”
  - Word stemming
    - e.g. “computer”, “computing”, “computerize” => “compute”
- **How to assign weights**
  - Not all words are equally important: Some are more indicative than others
    - e.g. “algebra” vs. “science”
- How to measure the similarity

# How to Assign Weights

- Two-fold heuristics based on frequency
  - TF (Term frequency)
    - More frequent **within** a document → more relevant to semantics
    - e.g., “query” vs. “commercial”
  - IDF (Inverse document frequency)
    - Less frequent **among** documents → more discriminative
    - e.g. “algebra” vs. “science”

# Term Weights: Term Frequency

- More frequent terms in a document are more important, i.e. more indicative of the topic.

$f_{ij}$  = frequency of term  $i$  in document  $j$

- May want to normalize *term frequency* ( $tf$ ) by dividing by the frequency of the most common term in the document:

$$tf_{ij} = f_{ij} / \max_i \{f_{ij}\}$$

# Term Weights: Inverse Document Frequency

- Terms that appear in many *different* documents are *less* indicative of overall topic.

$df_i$  = document frequency of term  $i$

= number of documents containing term  $i$

$idf_i$  = inverse document frequency of term  $i$ ,

=  $\log_{10} (N / df_i)$

( $N$ : total number of documents)

- An indication of a term's *discrimination* power.
- Log used to dampen the effect relative to  $tf$ .

idf example, suppose  $N = 1$  million

term	$df_t$	$idf_t$
calpurnia	1	6
animal	100	
sunday	1,000	
fly	10,000	
under	100,000	
the	1,000,000	0

$$idf_t = \log_{10} (N/df_t)$$

There is one idf value for each term  $t$  in a collection.



# TF-IDF Weighting

- A typical combined term importance indicator is *tf-idf weighting*:

$$w_{ij} = tf_{ij} idf_i = tf_{ij} \log_{10} (N / df_i)$$

- A term occurring frequently in the document but **rarely in the rest of the collection** is given **high weight**.
- Many other ways of determining term weights have been proposed.
- Experimentally, *tf-idf* has been found to work well.



## Binary $\rightarrow$ count $\rightarrow$ weight matrix

---

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	5.25	3.18	0	0	0	0.35
Brutus	1.21	6.1	0	1	0	0
Caesar	8.59	2.54	0	1.51	0.25	0
Calpurnia	0	1.54	0	0	0	0
Cleopatra	2.85	0	0	0	0	0
mercy	1.51	0	1.9	0.12	5.25	0.88
worser	1.37	0	0.11	4.15	0.25	1.95

Each document is now represented by a real-valued vector of tf-idf weights  $\in \mathbb{R}^{|V|}$

# How to Measure Similarity?

- Given two document

$$D_i = (w_{i1}, w_{i2}, \dots, w_{iN})$$

$$D_j = (w_{j1}, w_{j2}, \dots, w_{jN})$$

- Similarity definition

- dot product

$$Sim(D_i, D_j) = \sum_{t=1}^N w_{it} * w_{jt}$$

- normalized dot product (or cosine)

$$Sim(D_i, D_j) = \frac{\sum_{t=1}^N w_{it} * w_{jt}}{\sqrt{\sum_{t=1}^N (w_{it})^2 * \sum_{t=1}^N (w_{jt})^2}}$$

# Illustrative Example

